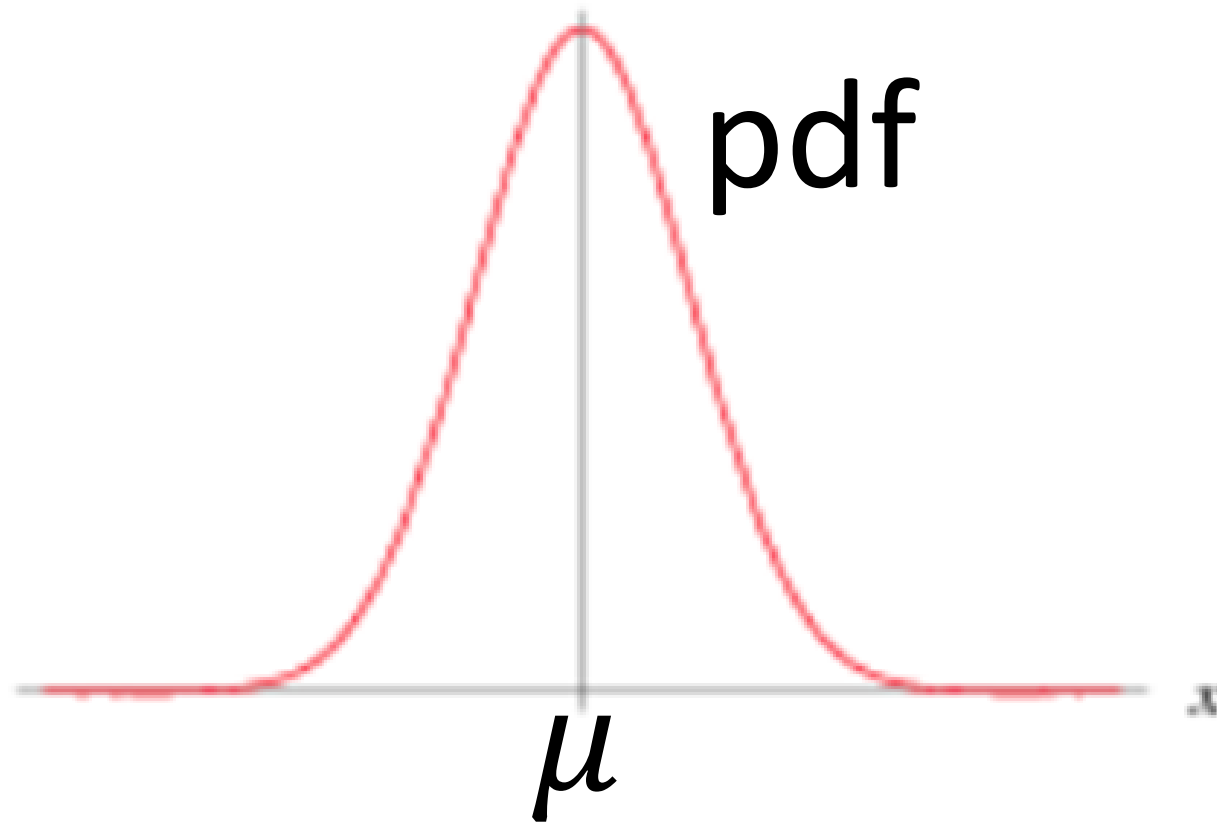


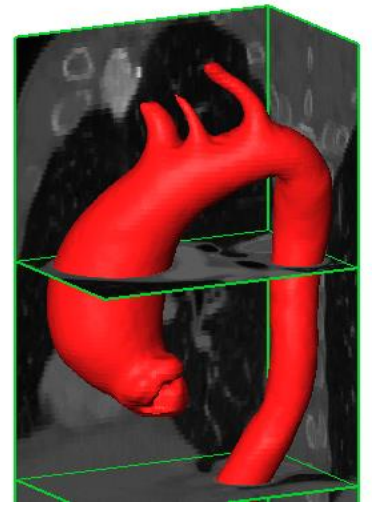
# 1-dimensional (1D) Gaussian Distribution

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Normal Distribution



$\{x_1, x_2, x_3, \dots, x_N\}$  is a set of data samples collected from a hospital  
 $x_n$  is the aorta size of a patient



Patient Population  
aorta size : 2 ~ 6 cm  
aorta size distribution



$x_1$  is the aorta size of the patient #1  
 $x_2$  is the aorta size of the patient #2  
 $x_3$  is the aorta size of the patient #3  
...

There could be nothing random about the aorta size and aneurysm of a patient

- the disease process could be deterministic

The data distribution can be model by a probability distribution (e.g. Gaussian)

In order to use the data distribution, we treat aorta size as random variable

disease group (ruptured)

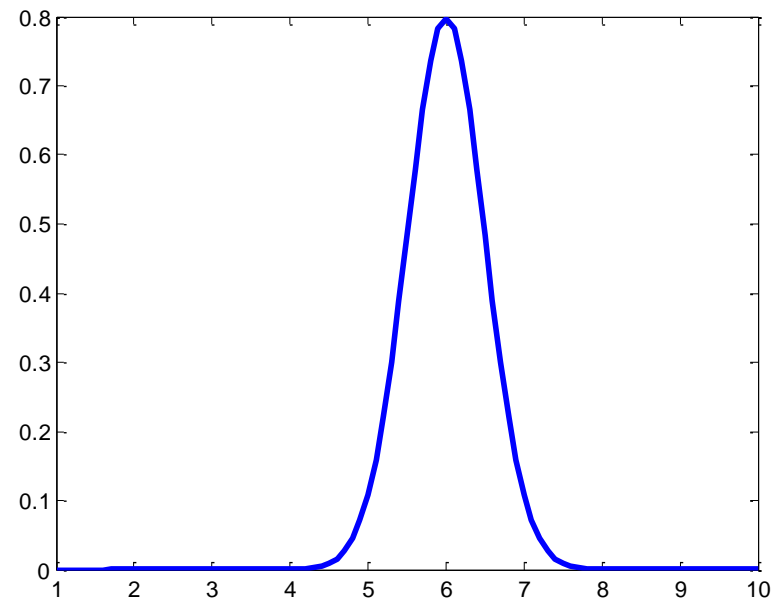
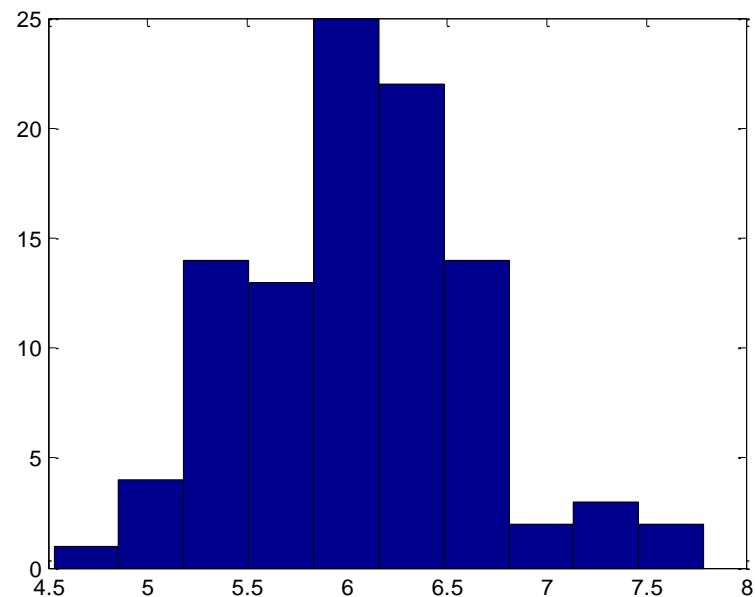
subject id	diameter
1	5.5
2	4.6
3	3.2
4	4.2
5	6.0

After we check the data, we guess that it may has a Gaussian distribution.

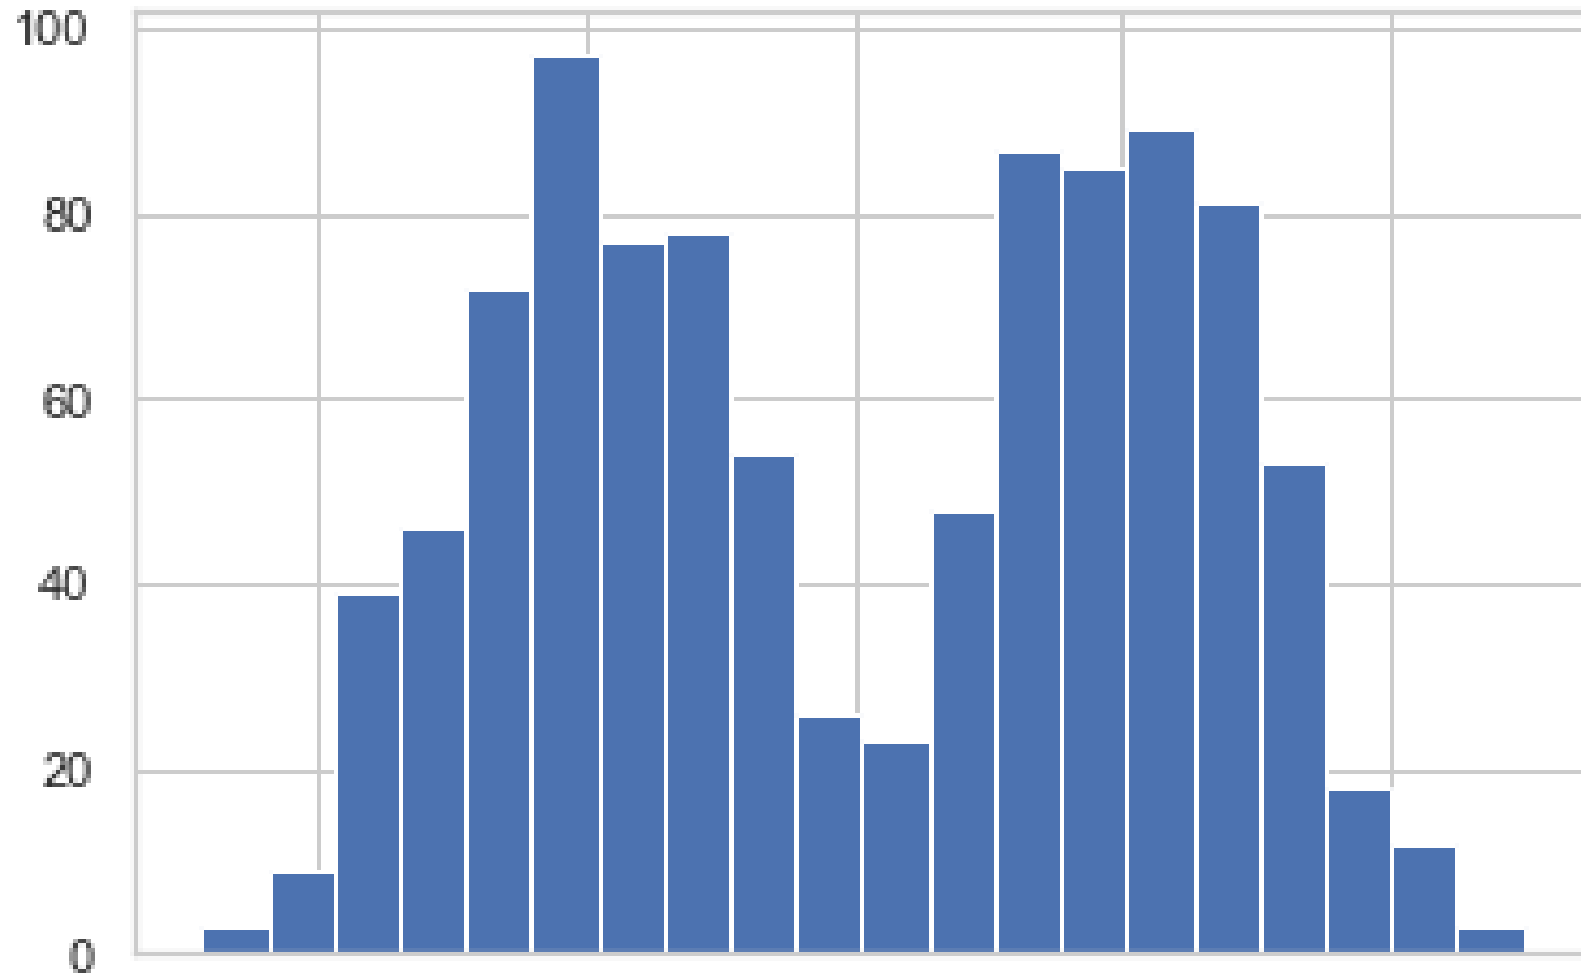
We use Gaussian PDF as the model of data distribution for machine learning.

Gaussian PDF

assume we have 100 subjects in this group



# What if the distribution is not a simple Gaussian ?



# Gaussian Mixture Model (GMM)

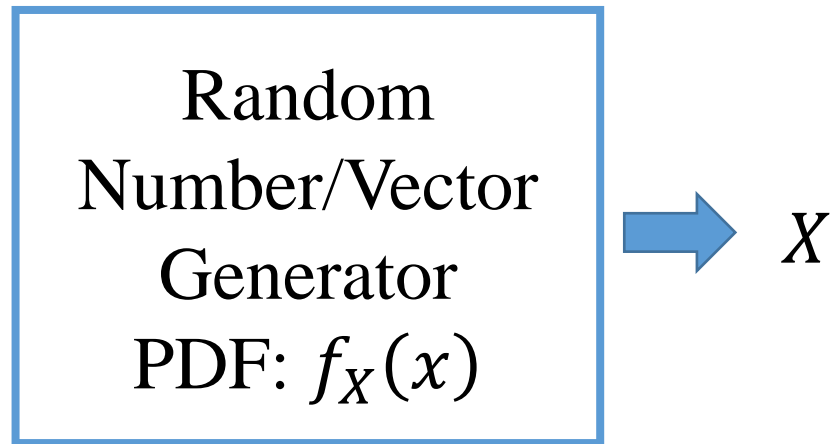
Liang Liang

## What does this mean?

$\{x_1, x_2, x_3, \dots, x_N\}$  is a set of data samples, and  $x_n \in \mathcal{R}^M$

$x_n$  is an observed value /realization of a random variable  $X_n$

random variables  $\{X_1, X_2, X_3, \dots, X_N\}$  are i.i.d. (**i**ndependently and **i**dentically **d**istributed)



$x_1$  is an observed value of  $X$   
 $x_2$  is an observed value of  $X$   
 $x_3$  is an observed value of  $X$   
we say  $x_1, x_2, x_3, \dots$  are  
generated from  $f_X(x)$

*Where are those r.v. ?*  
 $\{X_1, X_2, X_3, \dots, X_N\}$

```
In [1]: 1 import numpy as np
```

```
In [2]: 1 rng = np.random.RandomState()
```

```
In [3]: 1 x1 = rng.randn()  
2 x1
```

```
Out[3]: 0.5545141770294396
```

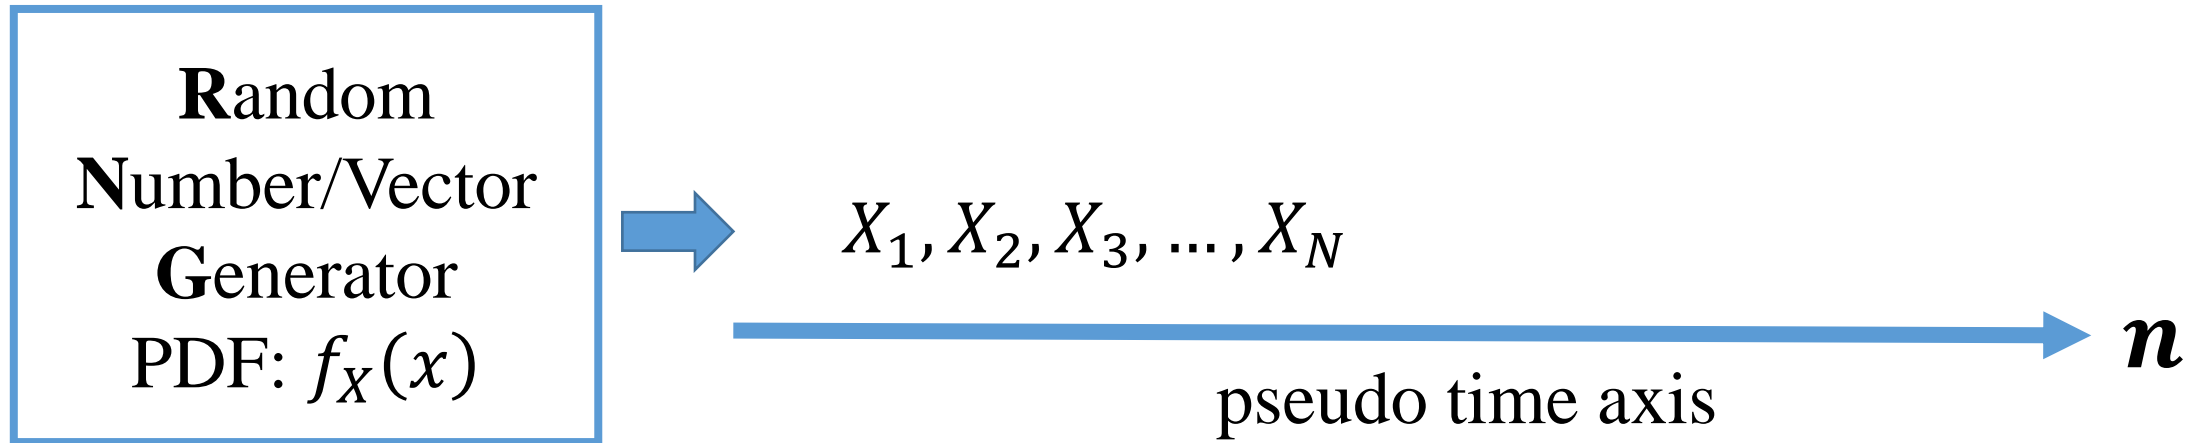
```
In [4]: 1 x2 = rng.randn()  
2 x2
```

```
Out[4]: -0.7736950982665748
```

```
In [5]: 1 x3 = rng.randn()  
2 x3
```

```
Out[5]: 0.5774013368938337
```

# Re-think the process from a different perspective



Task: Use the **rng** to generate a sequence of random numbers

After running Python code, we see a sequence of numbers coming out of **rng**

x1 is an observed value of  $X_1$

x2 is an observed value of  $X_2$

x3 is an observed value of  $X_3$

$\{X_1, X_2, X_3\}$  are **i.i.d.** because they come from the same PDF  
and they are generated independently



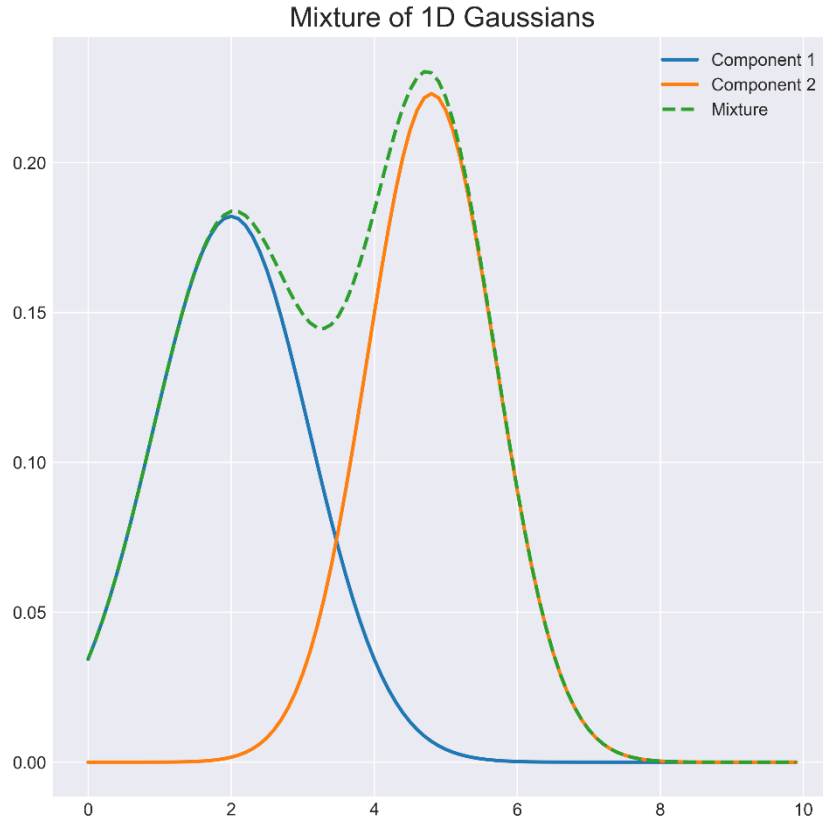
# Generative Model

- Assumption in machine learning:  
data samples are generated from a probability distribution
- A model that can generate data samples, is a generative model
- GMM is a probability distribution and a generative model

# 1D Gaussian Mixture Model

- Mixture of two Gaussians in 1D

$$f_X(x) = \pi_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} + \pi_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}}$$



There are two components/clusters

a r.v.  $X$  may belong to `cluster_1` or `cluster_2`

$\pi_1 = P(\{X \in cluster_1\})$  prior probability

$\pi_2 = P(\{X \in cluster_2\})$  prior probability

$\pi_1 + \pi_2 = 1$

$f_X(x)$  is the PDF of the random variable  $X$

cluster-1:  $\mathcal{N}(x|\mu_1, \sigma_1^2) = p_X(x|X \in cluster_1)$

cluster-2:  $\mathcal{N}(x|\mu_2, \sigma_2^2) = p_X(x|X \in cluster_2)$

See the demo: 1D\_GMM.ipynb

# Estimate the parameters of 1D GMM

- GMM with two 1D Gaussian components, the PDF is

$$f_X(x) = \pi_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} + \pi_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}}$$

- The GMM (PDF) has six **parameters**:  $\pi_1, \mu_1, \sigma_1, \pi_2, \mu_2, \sigma_2$
- Given a set of data samples, we fit the GMM model to the data.
- **Model fitting/training** is to estimate/recover the model parameters

# Maximum Likelihood (ML) Estimation

- $\{x_1, x_2, x_3, \dots, x_N\}$  is a set of data samples, and  $x_n \in \mathcal{R}^M$
- Assume  $x_n$  is an observed value of a random variable  $X_n$ , and the r.v.  $\{X_1, X_2, X_3, \dots, X_N\}$  are i.i.d. (independently and identically distributed)
- $f_{X_n}(x)$  is the PDF of  $X_n$  i.i.d.  $\Rightarrow f(x) = f_{X_1}(x) = \dots = f_{X_N}(x)$
- $f(x)$  has some parameters but the parameter values are unknown
- The goal is to estimate the parameters of  $f(x)$  from the data samples

# Maximum Likelihood (ML) Estimation

- The joint PDF is

$$f_{X_1, X_2, \dots, X_N}(x_1, x_2, \dots, x_N) = f_{X_1}(x_1) f_{X_2}(x_2) \dots f_{X_N}(x_N) = \prod_{n=1}^N f(x_n)$$

- The **NLL loss function** of the parameters is

$$loss = -\log f_{X_1, \dots, X_N}(x_1, \dots, x_N) = -\sum_{n=1}^N \log f(x_n)$$

# Maximum Likelihood (ML) Estimation

- The loss function of the parameters is

$$loss = -\log f_{X_1, \dots, X_N}(x_1, \dots, x_N) = -\sum_{n=1}^N \log f(x_n)$$

- $x_1, \dots, x_N$  are data points
- $-\log f_{X_1, \dots, X_N}(x_1, \dots, x_N)$  is the **Negative Log Likelihood (NLL)**
- The goal is to maximize the log Likelihood (to minimize NLL) to obtain the optimal parameters of  $f(x)$
- Why Max Likelihood ?

We want to find the best parameters such that the data samples are highly likely coming/generated from the PDF  $f(x)$

The observed data samples are most probable under the assumed PDF

# Maximum Likelihood (ML) Estimation of 1D Gaussian

- Let's assume the PDF is a 1D Gaussian

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

The NLL loss function is

$$loss = L(\mu, \sigma) = -\sum_{n=1}^N \log f(x_n) = \sum_{n=1}^N \frac{(x_n - \mu)^2}{2\sigma^2} + N \log(\sqrt{2\pi}\sigma)$$

The loss reaches the minimum when  $\frac{\partial L}{\partial \mu} = 0$  and  $\frac{\partial L}{\partial \sigma} = 0$  (necessary condition)

From  $\frac{\partial L}{\partial \mu} = 0$ , we get  $\mu = \frac{1}{N} \sum_{n=1}^N x_n$

From  $\frac{\partial L}{\partial \sigma} = 0$ , we get  $\sigma^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu)^2$



# GMM for Clustering 1D data

- GMM with two Gaussian components, the PDF is

$$f_X(x) = \pi_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} + \pi_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}}$$

- Each component corresponds to a cluster
- After model fitting, we do **clustering**: assign each data point to a cluster
- For the data point  $x = x_n$ ,

if  $\pi_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} > \pi_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}}$ , then  $x_n$  is assigned to cluster-1

if  $\pi_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} < \pi_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}}$ , then  $x_n$  is assigned to cluster-2

# M-dimensional (M-D) Gaussian Distribution

$$\bullet f_X(x) = \frac{1}{(2\pi)^{M/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

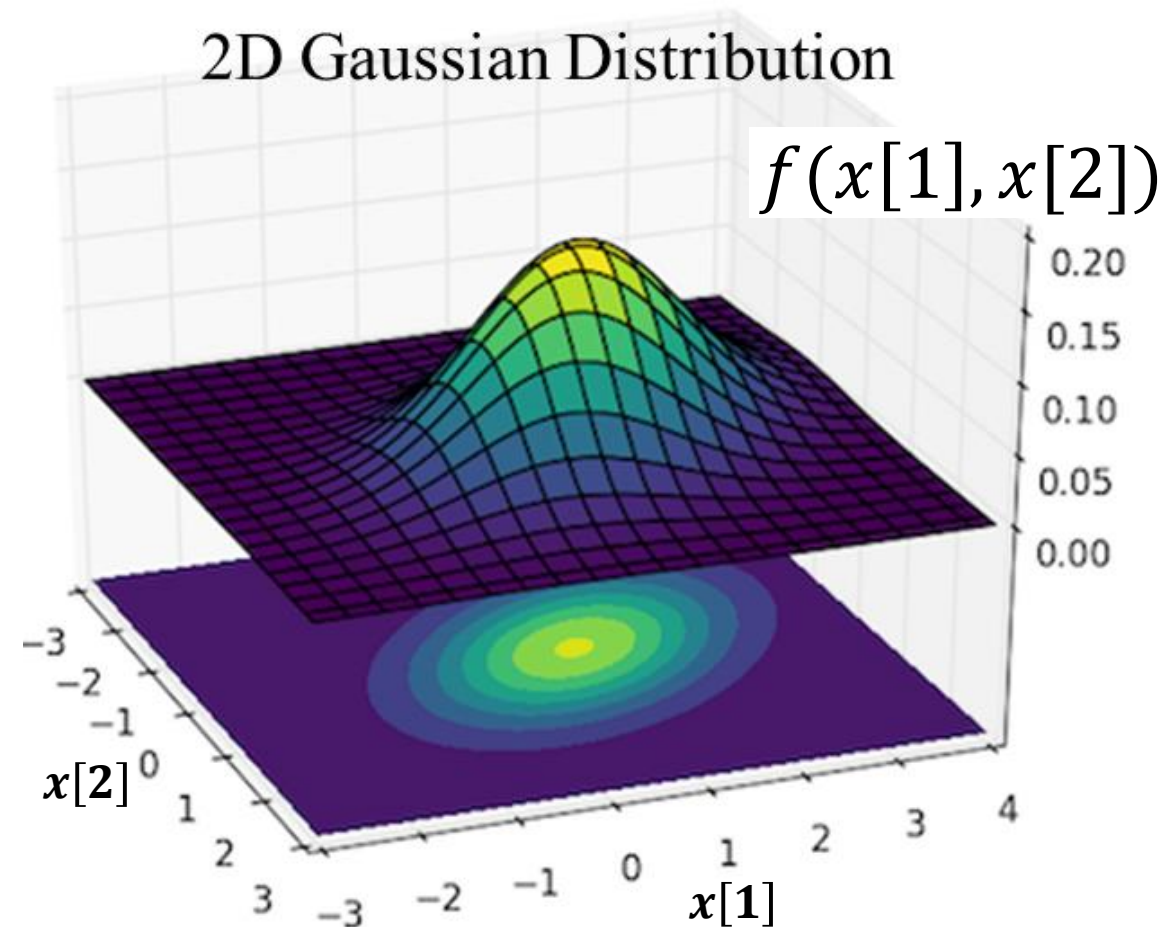
$$x, \mu \in \mathcal{R}^M, \Sigma \in \mathcal{R}^{M \times M}$$

$$f_X(x)$$

$$= \mathcal{N}(\mu, \Sigma)$$

$$= \mathcal{N}(x; \mu, \Sigma)$$

$$= \mathcal{N}(x | \mu, \Sigma)$$



# M-D Gaussian Mixture Model

- Mixture of  $K$  Gaussians in  $M$ -dim:  $x, \mu \in \mathcal{R}^M, \Sigma \in \mathcal{R}^{M \times M}$

$$f_X(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \mu_k, \Sigma_k)$$

$$\pi_k = P(\{X \in cluster_k\}), \text{ and } \sum_{k=1}^K \pi_k = 1$$

- the GMM has  $K$  Gaussian components/clusters.
- the GMM is the PDF of a random variable  $X$
- a data sample  $x$  is an observed value of the random variable  $X$

# Maximum Likelihood Estimation of GMM Parameters

- The loss function of the parameters is

$$loss = -\log f_{X_1, \dots, X_N}(x_1, \dots, x_N) = -\sum_{n=1}^N \log f(x_n)$$

- $x_1, \dots, x_N$  are data points
- $-\log f_{X_1, \dots, X_N}(x_1, \dots, x_N)$  is **Negative Log Likelihood (NLL)**
- The goal is to maximize the log Likelihood (to minimize NLL) to obtain the optimal parameters of the PDF  $f(x)$
- GMM parameters:  $\pi_k, \mu_k, \Sigma_k$  to be estimated from data samples

# Maximum Likelihood Estimation of GMM Parameters

- **GMM:**  $f(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$  where  $\sum_{k=1}^K \pi_k = 1$
- **Goal:** estimate  $\mu_k, \Sigma_k, \pi_k$  for  $k=1$  to  $K$ , from data samples
- **Define:**  $\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)}$  and  $N_k = \sum_{n=1}^N \gamma_{(n,k)}$

$$\gamma_{(n,k)} = P(\{X_n \in cluster_k\} | X_n = x_n)$$

$\gamma_{(n,k)}$  is called the posterior probability

It is the probability of  $X_n$  belonging to  $cluster_k$  after we observe its value  $x_n$ . In layman's term, it is the probability of  $x_n$  belonging to the  $cluster_k$ .

# Maximum Likelihood Estimation of GMM Parameters

- $\gamma_{(n,k)} = P(\{X_n \in cluster_k\} | X_n = x_n)$ 
  - $\gamma_{(n,k)}$  is the probability of sample  $x_n$  belonging to  $cluster_k$
  - $\gamma_{(n,k)}$  is the posterior probability, soft assignment/label
- Define  $z_{(n,k)}$  to be the indicator of the event  $\{X_n \in cluster_k\}$ 
  - If  $z_{(n,k)} = 1$ , then it is true that  $X_n \in cluster_k$
  - If  $z_{(n,k)} = 0$ , then it is false that  $X_n \in cluster_k$
  - $z_{(n,k)}$  is the (hard) cluster label/indicator of  $x_n$

# Maximum Likelihood Estimation of GMM Parameters

$$\gamma_{(n,k)} = P(\{X_n \in cluster_k\} | X_n = x_n)$$

$z_{(n,k)}$  is the binary indicator of the event  $\{X_n \in cluster_k\}$

$$\gamma_{(n,k)} = p(z_{(n,k)} = 1 | X_n = x_n)$$

Let's simplify the notations:

$$\gamma_{(n,k)} = p(z_{(n,k)} | x_n)$$

It is the probability of  $x_n$  belonging to cluster- $k$

# Maximum Likelihood Estimation of GMM Parameters

- GMM:  $f(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$  where  $\sum_{k=1}^K \pi_k = 1$

$$\gamma_{(n,k)} = p(z_{(n,k)}|x_n) = \frac{p(z_{(n,k)}, x_n)}{p(x_n)} = \frac{p(z_{(n,k)}, x_n)}{f(x_n)}$$

**assume  $f(x)$  is the true PDF**

$$p(z_{(n,k)}, x_n) = p(z_{(n,k)})p(x_n|z_{(n,k)})$$

$$p(z_{(n,k)}) \triangleq \pi_k, \text{ and } p(x_n|z_{(n,k)}) \triangleq \mathcal{N}(x_n|\mu_k, \Sigma_k)$$

$$\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)}$$



# Maximum Likelihood (ML) Estimation of GMM Parameters

- The NLL loss function of the parameters is

$$L = - \sum_{n=1}^N \log f(x_n)$$

- GMM:  $f(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$  where  $\sum_{k=1}^K \pi_k = 1$

- To obtain  $\mu_k$ :  $\frac{\partial L}{\partial \mu_k} = 0$ , then we have

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} x_n$$

where  $N_k = \sum_{n=1}^N \gamma_{(n,k)}$

# Maximum Likelihood Estimation of GMM Parameters

- The ML loss (NLL) function of the parameters is

$$L = - \sum_{n=1}^N \log f(x_n)$$

- GMM:  $f(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$  where  $\sum_{k=1}^K \pi_k = 1$

- To obtain  $\Sigma_k$ :  $\frac{\partial L}{\partial \Sigma_k} = 0$ , then we have

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} (x_n - \mu_k)(x_n - \mu_k)^T$$

$\Sigma_k$  is called covariance matrix

where  $N_k = \sum_{n=1}^N \gamma_{(n,k)}$

# Maximum Likelihood Estimation of GMM Parameters

- The ML loss (NLL) function of the parameters is

$$L = - \sum_{n=1}^N \log f(x_n)$$

- GMM:  $f(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$  where  $\sum_{k=1}^K \pi_k = 1$
- To obtain  $\pi_k$ : we add a new term to  $L$  *because*  $\sum_{k=1}^K \pi_k = 1$

$$L = - \sum_{n=1}^N \log f(x_n) + \lambda \left( \sum_{k=1}^K \pi_k - 1 \right)$$

$$\frac{\partial L}{\partial \pi_k} = 0, \text{ then we have } \pi_k = \frac{N_k}{N}$$

$$\text{where } N_k = \sum_{n=1}^N \gamma_{(n,k)}$$

# Algorithm for Model Fitting to Estimate Parameters

- **Input:**  $\{x_1, x_2, x_3, \dots, x_N\}$  is a set of data samples, and  $x_n \in \mathcal{R}^M$ ,  
 $K$  is the number of components
- **Initialization:** initial values of  $\mu_k, \Sigma_k, \pi_k$  for  $k=1$  to  $K$
- **E-step:** compute the soft assignment  $\gamma_{(n,k)}$ , the probability of sample  $x_n$  belonging to *cluster* <sub>$k$</sub>

$$\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$

soft assignment matrix  $[\gamma_{(n,k)}]$

- **M-step:** update the parameters  $\mu_k, \Sigma_k, \pi_k$  for  $k=1$  to  $K$

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} x_n$$

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} (x_n - \mu_k)(x_n - \mu_k)^T$$

$$\pi_k = \frac{N_k}{N}, \text{ where } N_k = \sum_{n=1}^N \gamma_{(n,k)}$$

- Iterate the two steps until convergence: loss almost does not change

# Clustering is Performed After Model Fitting

- **Input:**  $\{x_1, x_2, x_3, \dots, x_N\}$  is a set of data samples, and  $x_n \in \mathcal{R}^M$ ,  
 $K$  is the number of components
- Run the Parameter Estimation Algorithm and we obtain the final soft assignment  $\gamma_{(n,k)}$ , the probability of sample  $x_n$  belonging to *cluster* <sub>$k$</sub>

$$\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)} \text{ where } \sum_{k=1}^K \gamma_{(n,k)} = 1$$

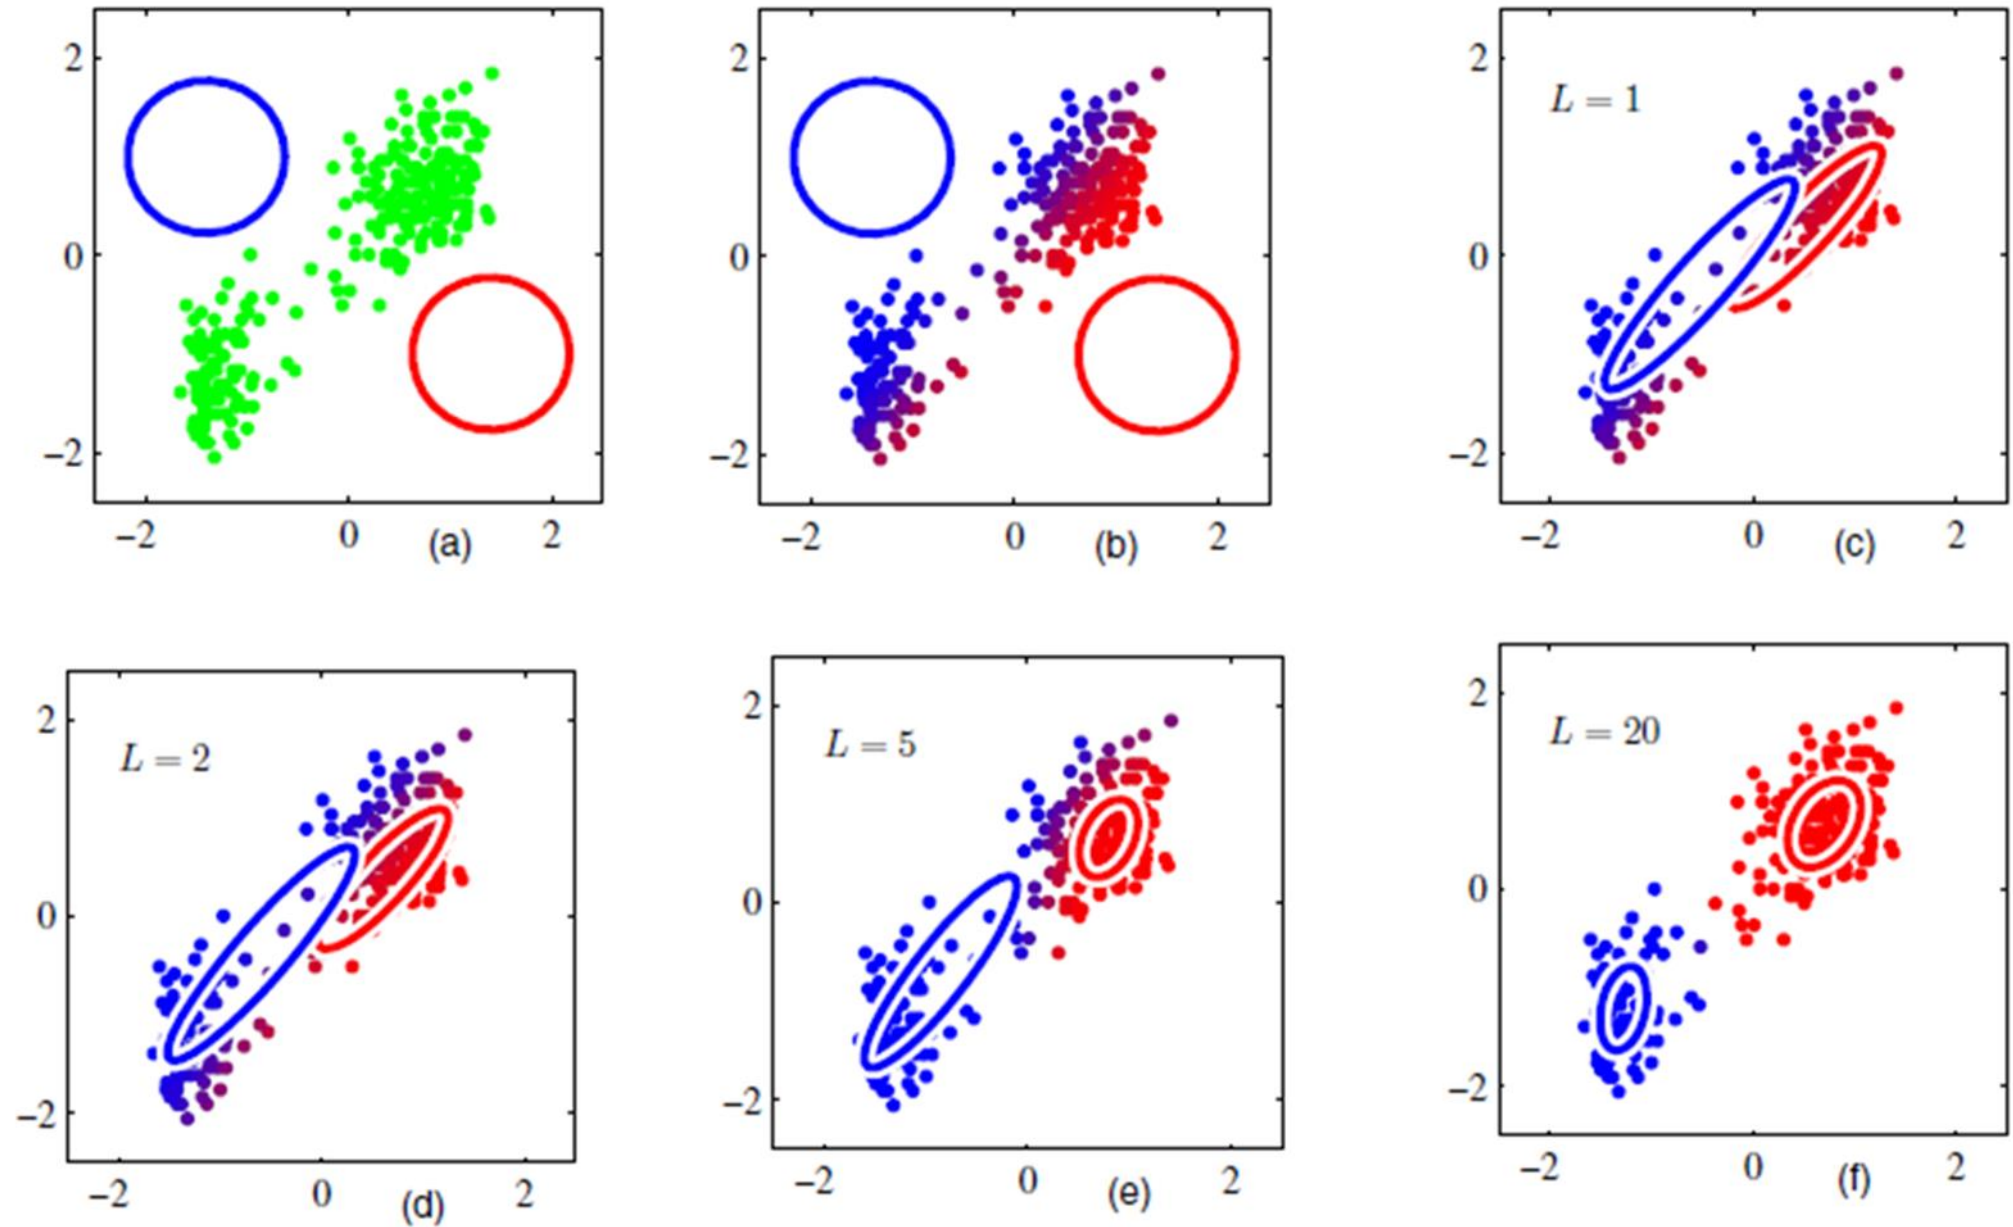
soft assignment matrix  $[\gamma_{(n,k)}]$

- **Clustering:**

cluster label of  $x_n$  is  $\operatorname{argmax}[\gamma_{(n,1)}, \gamma_{(n,2)}, \gamma_{(n,3)}, \dots, \gamma_{(n,K)}]$

example:  $x_n$  is assigned to cluster-1 if the row- $n$  is  $[0.9, 0.1, 0.01, \dots]$

# Initialization



# EM Algorithm

- The algorithm is derived by setting  $\frac{\partial L}{\partial \mu_k} = 0$ ,  $\frac{\partial L}{\partial \Sigma_k} = 0$  and  $\frac{\partial L}{\partial \pi_k} = 0$
- The algorithm can also be derived from a method: **Expectation-Maximization**
  - E-step: calculate  $f^t(\mathbf{J}) \triangleq P(\mathbf{J}|\mathbf{U}, \Theta^t)$
  - M-step:  $\Theta^{t+1} = \operatorname{argmax}_{\Theta} [Q^t(\Theta) + \log P(\Theta)]$

It is very useful if parameter estimation is a 'chicken-egg' problem. A model has two sets of parameters, set-a, set-b.

(1) to get set-a, we need to know set-b

(2) to get set-b, we need to know set-a

## The Expectation Maximization Algorithm

Frank Dellaert

College of Computing, Georgia Institute of Technology

Technical Report number GIT-GVU-02-20

February 2002

### Abstract

This note represents my attempt at explaining the EM algorithm (Hartley, 1958; Dempster et al., 1977; McLachlan and Krishnan, 1997). This is just a slight variation on Tom Minka's tutorial (Minka, 1998), perhaps a little easier (or perhaps not). It includes a graphical example to provide some intuition.

# How can we determine the optimal value of $K$ ?

- Knowledge of the application
- Bayesian information criterion (BIC)
- Akaike information criterion (AIC)
- The basic idea is to reduce the number of components while keeping a high log likelihood. For example, you may define your own criterion:

$$\text{loss}(K) = -\text{Log-likelihood}/N + g(K)$$



# GMM => modified k-means for clustering

- **E-step (GMM):** compute the soft assignment  $\gamma_{(n,k)}$ , the probability of sample  $x_n$  belonging to *cluster* <sub>$k$</sub>

$$\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)} \text{ where } \sum_{k=1}^K \gamma_{(n,k)} = 1$$

- **E-step (modified k-means):** compute the soft assignment  $\gamma_{(n,k)}$ , and convert it to hard assignment.

e.g.  $K=2$  soft  $\gamma_{(n,1)} = 0.1, \gamma_{(n,2)} = 0.9 \Rightarrow$  hard  $\gamma_{(n,1)} = 0, \gamma_{(n,2)} = 1$

- **M-step:** update the parameters  $\mu_k, \Sigma_k, \pi_k$  for  $k=1$  to  $K$

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} x_n$$

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_{(n,k)} (x_n - \mu_k)(x_n - \mu_k)^T$$

$$\pi_k = \frac{N_k}{N}, \text{ where } N_k = \sum_{n=1}^N \gamma_{(n,k)}$$

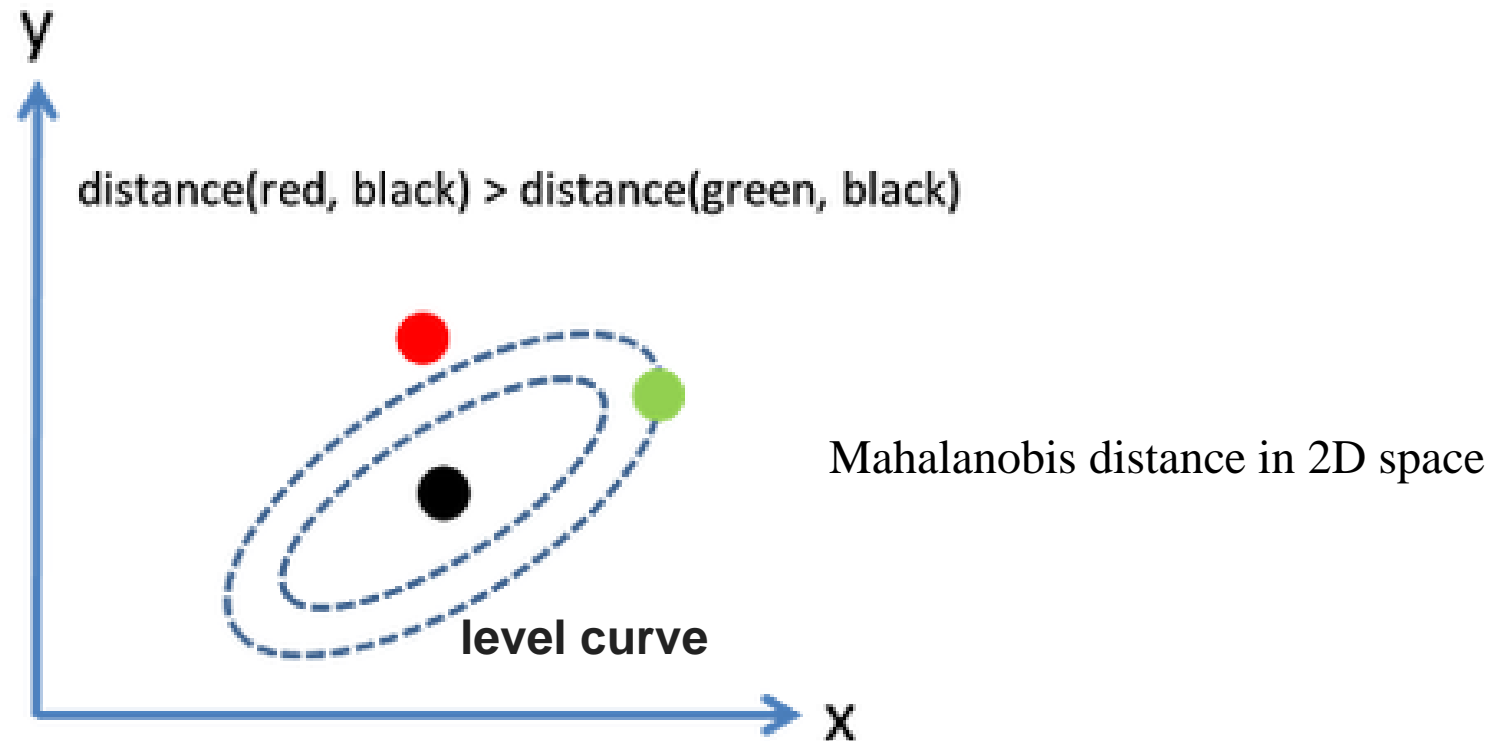
- Iterate the two steps until convergence: loss almost does not change

# The distance function of the modified k-means

- The distance function is

$$d_k(x_n, \mu_k) = \sqrt{(x_n - \mu_k)^T \Sigma_k^{-1} (x_n - \mu_k)}$$

**a.k.a.** Mahalanobis distance



# The distance function of the modified k-means

- **E-step (modified k-means):** compute the soft assignment  $\gamma_{(n,k)}$ , and convert it to hard assignment.

$$\gamma_{(n,k)} = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)} \text{ where } \sum_{k=1}^K \gamma_{(n,k)} = 1$$

set  $K=2$ :

$$\gamma_{(n,1)} = \frac{\pi_1 \mathcal{N}(x_n | \mu_1, \Sigma_1)}{\sum_{j=1}^2 \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}, \quad \gamma_{(n,2)} = \frac{\pi_2 \mathcal{N}(x_n | \mu_2, \Sigma_2)}{\sum_{j=1}^2 \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)} \text{ then } \gamma_{(n,1)} + \gamma_{(n,2)} = 1$$

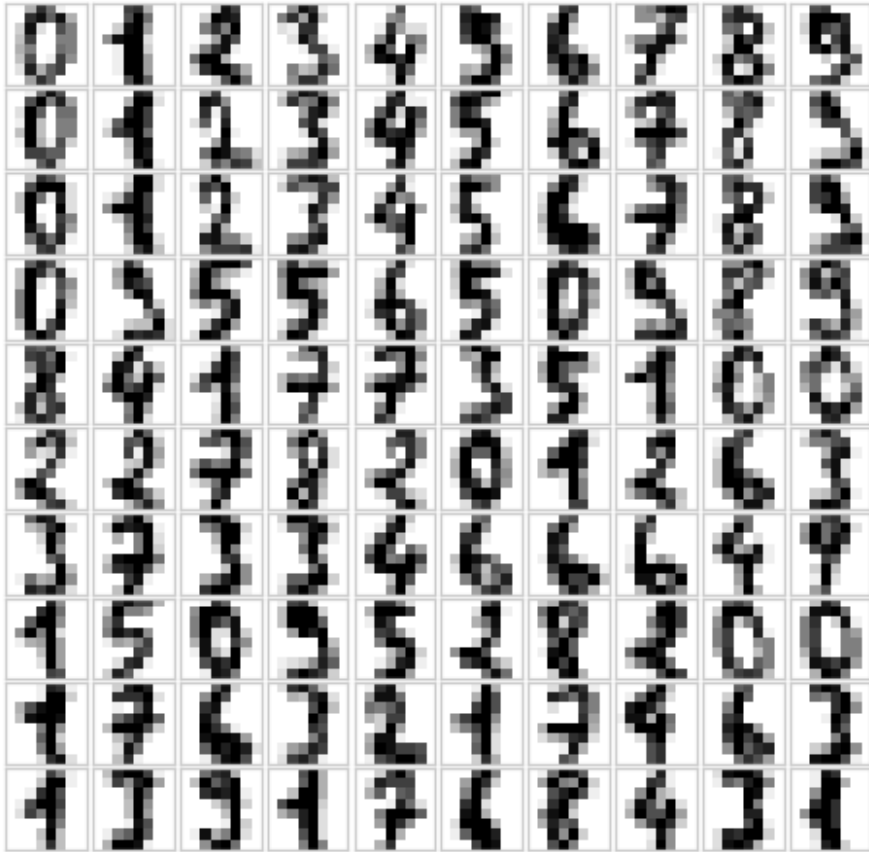
if  $\gamma_{(n,1)} = 0$ , then  $x_n$  belongs to cluster-2

else: ratio =  $\frac{\gamma_{(n,2)}}{\gamma_{(n,1)}} = \frac{\pi_2 \mathcal{N}(x_n | \mu_2, \Sigma_2)}{\pi_1 \mathcal{N}(x_n | \mu_1, \Sigma_1)}$ , if ratio  $> 1$ , then  $x_n \in cluster_2$

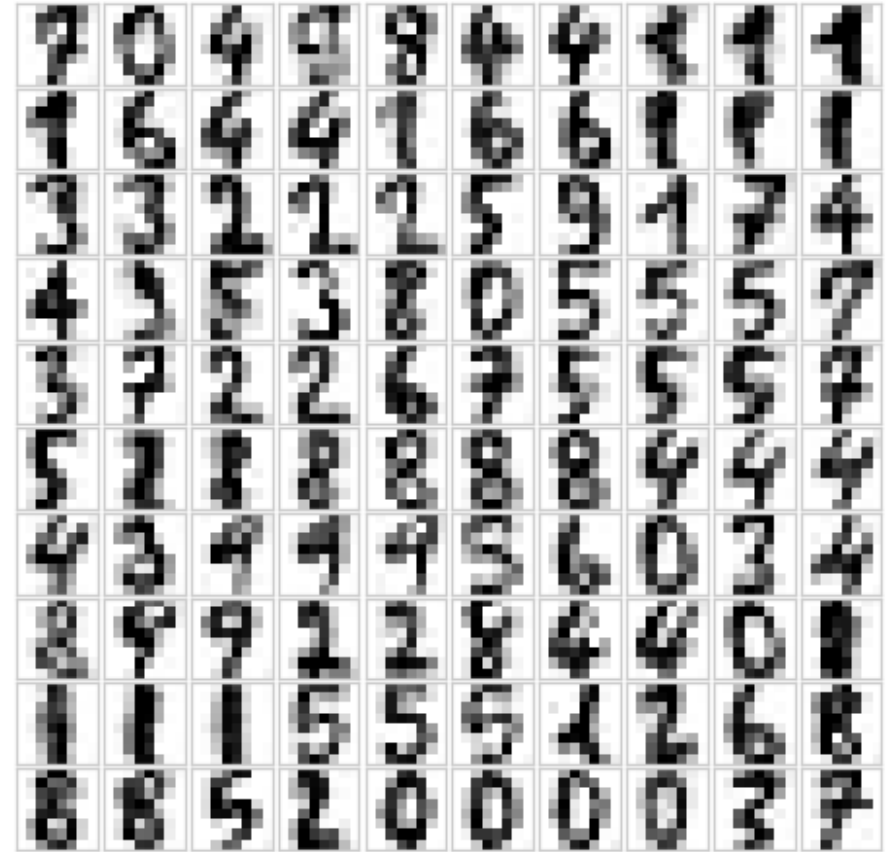
$$\begin{aligned} \log \frac{\gamma_{(n,2)}}{\gamma_{(n,1)}} &\approx (x_n - \mu_1)^T \Sigma_1^{-1} (x_n - \mu_1) - (x_n - \mu_2)^T \Sigma_2^{-1} (x_n - \mu_2) \\ &= (d_1(x_n, \mu_1))^2 - (d_2(x_n, \mu_2))^2 \end{aligned}$$

GMM is a Generative Model because it is a PDF

Training Data



Generated New Data



see demos:

2D\_GMM.ipynb

GMM\_DE\_Generative\_Model.ipynb